

In []: ▶

```

1  #####
2  #
3  #               Project Assignment on Convolution neural network
4  #
5  #####
6  # Problem statement: To Classify the images from CIFAR-10 DataSet Using CNN
7
8
9  # Loading the Data ,by unzipping the tar file and creating batches of images
10 # Install tqdm file
11
12 from urllib.request import urlretrieve
13 from os.path import isfile, isdir
14
15 from tqdm import tqdm
16 import tarfile
17
18
19 cifar10_dataset_folder_path = 'cifar-10-batches-py'
20
21 class DLProgress(tqdm):
22     last_block = 0
23
24     def hook(self, block_num=1, block_size=1, total_size=None):
25         self.total = total_size
26         self.update((block_num - self.last_block) * block_size)
27         self.last_block = block_num
28
29 # if the file cifar tar is not downloaded , then download
30 if not isfile('cifar-10-python.tar.gz'):
31     with DLProgress(unit='B', unit_scale=True, miniters=1, desc='CIFAR-10 Dataset'):
32         urlretrieve('https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz',
33                    'cifar-10-python.tar.gz', pbar.hook)
34
35 # open the file and extract all the batches from the tar file
36 if not isdir(cifar10_dataset_folder_path):
37     with tarfile.open('cifar-10-python.tar.gz') as tar:
38         tar.extractall()
39         tar.close()

```

```

In [ ]: ▶ 1 #Each batch had many images , so there are 4 dimension , first value is images , 2nd and 3rd size and colour
2 # in transpose we mention the index number of the dimensions
3 # The input features are normalised to save on memory , for accuracy and within a particular scale
4 # pooling can be differnet types, one of them is max pooling
5
6 import pickle
7 #pickle is for serializing and unserializing the image data into bits/bytes
8 import numpy as np
9 import matplotlib.pyplot as plt
10 from sklearn.preprocessing import LabelBinarizer
11
12
13 # Defining function to Load the names of the Labels of whose images in the DataSet
14 def load_label_names():
15     """    Load the label names from file    """
16     return ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
17
18
19 # Function to convert the input image data into the way tensor for Convolution accepts i.e.width X height X chan
20 # and returns the features and Labels containing the data from the batch
21
22 def load_cifar10_batch(cifar10_dataset_folder_path, batch_id):
23     """    Load a batch of the dataset    """
24     with open(cifar10_dataset_folder_path + '/data_batch_' + str(batch_id), mode='rb') as file:
25         # note the encoding type is 'latin1'
26         batch = pickle.load(file, encoding='latin1')
27
28         features = batch['data'].reshape((len(batch['data']), 3, 32, 32)).transpose(0, 2, 3, 1)
29         labels = batch['labels']
30
31         return features, labels
32
33
34 # Function to show the image selected . It internally loads the image from the batch mentioned in
35 # the path given then reshapes the image
36
37 def display_stats(cifar10_dataset_folder_path, batch_id, sample_id):
38     """    Display Stats of the the dataset    """
39     batch_ids = list(range(1, 6))
40     if batch_id not in batch_ids:
41         print('Batch Id out of Range. Possible Batch Ids: {}'.format(batch_ids))

```

```
42     return None
43
44     features, labels = load_cifar10_batch(cifar10_dataset_folder_path, batch_id)
45
46     if not (0 <= sample_id < len(features)):
47         print('{} samples in batch {}. {} is out of range.'.format(len(features), batch_id, sample_id))
48         return None
49
50     print('\nStats of batch #{}:'.format(batch_id))
51     print('# of Samples: {}'.format(len(features)))
52
53     label_names = load_label_names()
54     label_counts = dict(zip(*np.unique(labels, return_counts=True)))
55     for key, value in label_counts.items():
56         print('Label Counts of [{}]({}) : {}'.format(key, label_names[key].upper(), value))
57
58     sample_image = features[sample_id]
59     sample_label = labels[sample_id]
60
61     print('\nExample of Image {}:".format(sample_id))
62     print('Image - Min Value: {} Max Value: {}'.format(sample_image.min(), sample_image.max()))
63     print('Image - Shape: {}'.format(sample_image.shape))
64     print('Label - Label Id: {} Name: {}'.format(sample_label, label_names[sample_label]))
65
66     plt.imshow(sample_image)
67
```

In []: ▶

```
1
2 %matplotlib inline
3 %config InlineBackend.figure_format = 'retina'
4
5 import numpy as np
6
7 # Explore the dataset
8 batch_id = 2
9 sample_id = 2700
10
11 cifar10_dataset_folder_path = "C:/Users/Maximus/pythonsessions/session 39Project_CNN/cifar-10-batches-py"
12 display_stats(cifar10_dataset_folder_path, batch_id, sample_id)
```

In []: ▶

```
1  # Normalisation : Function to normalize the Data , so that it lies within a range and we do not
2  # come across exploding gradient, here the value taken by input is in range of 0-255
3  # , this is also called squashing of the input value
4
5  from sklearn import preprocessing
6  lb = preprocessing.LabelBinarizer()
7  lb.fit(range(10))
8
9  def normalize(x):
10     """
11         argument
12         - x: input image data in numpy array [32, 32, 3]
13         return
14         - normalized x
15     """
16     # change 4
17     x = np.array(x)
18     min_val = np.min(x)
19     max_val = np.max(x)
20     x = (x-min_val) / (max_val-min_val)
21     return x
22
23 # One-Hot Encoding : Since we will be classifying the images as one of the output 10 categories
24 # and that is identified as highest probability of one of these 10 values. WE use one hot encoding
25 # to identify the output category
26 def one_hot_encode(x):
27     """
28         argument
29         - x: a list of labels
30         return
31         - one hot encoding matrix (number of labels, number of class)
32     """
33     return lb.transform(x)
34
35
```

```

In [ ]: 1 # Preprocess the Data , i.e normalize the input values of image and one hot encode the output
2 # values and write the new values to the file
3
4 def _preprocess_and_save(normalize, one_hot_encode, features, labels, filename):
5     features = normalize(features)
6     labels = one_hot_encode(labels)
7
8     pickle.dump((features, labels), open(filename, 'wb'))
9
10
11 # In this function , the image data for all the batches is preprocessed i.e normalized and one hot
12 # encoded. This data is then split into training and Validation Data , 10% of for validation
13 # Even Testing Data is preprocessed and dumped into the file
14 def preprocess_and_save_data(cifar10_dataset_folder_path, normalize, one_hot_encode):
15     n_batches = 5
16     valid_features = []
17     valid_labels = []
18
19     for batch_i in range(1, n_batches + 1):
20         features, labels = load_cifar10_batch(cifar10_dataset_folder_path, batch_i)
21
22         # find index to be the point as validation data in the whole dataset of the batch (10%)
23         index_of_validation = int(len(features) * 0.1)
24
25         # preprocess the 90% of the whole dataset of the batch
26         # - normalize the features
27         # - one_hot_encode the labels
28         # - save in a new file named, "preprocess_batch_" + batch_number
29         # - each file for each batch
30         _preprocess_and_save(normalize, one_hot_encode,
31                             features[:-index_of_validation], labels[:-index_of_validation],
32                             'preprocess_batch_' + str(batch_i) + '.p')
33
34         # unlike the training dataset, validation dataset will be added through all batch dataset
35         # - take 10% of the whole dataset of the batch
36         # - add them into a list of
37         #   - valid_features
38         #   - valid_labels
39         valid_features.extend(features[-index_of_validation:])
40         valid_labels.extend(labels[-index_of_validation:])
41

```

```
42 # preprocess the all stacked validation dataset
43 _preprocess_and_save(normalize, one_hot_encode,
44                     np.array(valid_features), np.array(valid_labels),
45                     'preprocess_validation.p')
46
47 # Load the test dataset
48 with open(cifar10_dataset_folder_path + '/test_batch', mode='rb') as file:
49     batch = pickle.load(file, encoding='latin1')
50
51 # preprocess the testing data
52 test_features = batch['data'].reshape((len(batch['data']), 3, 32, 32)).transpose(0, 2, 3, 1)
53 test_labels = batch['labels']
54
55 # Preprocess and Save all testing data
56 _preprocess_and_save(normalize, one_hot_encode,
57                     np.array(test_features), np.array(test_labels),
58                     'preprocess_training.p')
59
```

```
In [ ]: 1 # calling the function to preprocess the data and normalize and one hot encode the output
2 preprocess_and_save_data(cifar10_dataset_folder_path, normalize, one_hot_encode)
3
4
5 # Loading the data from the validation file
6 valid_features, valid_labels = pickle.load(open('preprocess_validation.p', mode='rb'))
```

```
In [ ]: 1 valid_features[:10] , valid_labels[:10]
```


In []: ▶

```

1
2 # Defining Convolution Layer
3
4 def conv2d_maxpool(x_tensor, conv_num_outputs, conv_ksize, conv_strides, pool_ksize, pool_strides):
5     """
6     Apply convolution then max pooling to x_tensor
7     :param x_tensor: TensorFlow Tensor
8     :param conv_num_outputs: Number of outputs for the convolutional layer
9     :param conv_ksize: kernel size 2-D Tuple for the convolutional layer
10    :param conv_strides: Stride 2-D Tuple for convolution
11    :param pool_ksize: kernel size 2-D Tuple for pool
12    :param pool_strides: Stride 2-D Tuple for pool
13    : return: A tensor that represents convolution and max pooling of x_tensor
14    """
15    # Weights
16    shape_of_x_tensor = x_tensor.get_shape().as_list()
17
18    F_W = tf.Variable(tf.truncated_normal([conv_ksize[0],conv_ksize[1],shape_of_x_tensor[-1],conv_num_outputs]
19                                         ,dtype=tf.float32,stddev=0.2))
20
21    F_b = tf.Variable(tf.zeros([conv_num_outputs],dtype=tf.float32))
22
23    strides_conv = [1,conv_strides[0],conv_strides[1],1]
24    padding = 'SAME'
25
26    output = tf.nn.conv2d(x_tensor, F_W, strides_conv, padding )
27    output = tf.nn.bias_add(output, F_b)
28
29    # Nonlinear activation (ReLU)
30    output = tf.nn.relu(output)
31
32    # Max pooling
33    ksize_maxpool = [1,pool_ksize[0],pool_ksize[1],1]
34    strides_maxpool = [1,pool_strides[0],pool_strides[1],1]
35    output = tf.nn.max_pool(output, ksize_maxpool ,strides_maxpool,padding)
36
37    return output
38

```

In []: ▶

```

1  #Defining the flatten layer , after the convolution layer
2
3  def flatten(x_tensor):
4      """
5      Flatten x_tensor to (Batch Size, Flattened Image Size)
6      : x_tensor: A tensor of size (Batch Size, ...), where ... are the image dimensions.
7      : return: A tensor of size (Batch Size, Flattened Image Size).
8      """
9      shape = x_tensor.get_shape().as_list()
10
11     flat_dim = shape[1] * shape[2] * shape[3]
12
13     output = tf.reshape(x_tensor, [-1, flat_dim])
14     return output
15
16 # defining the fully connected Layer
17
18 def fully_conn(x_tensor, num_outputs):
19     """
20     Apply a fully connected layer to x_tensor using weight and bias
21     : x_tensor: A 2-D tensor where the first dimension is batch size.
22     : num_outputs: The number of output that the new tensor should be.
23     : return: A 2-D tensor where the second dimension is num_outputs.
24     """
25     # Weights and bias
26     shape = x_tensor.get_shape().as_list()
27
28     W = tf.Variable(tf.truncated_normal([shape[1], num_outputs], dtype=tf.float32, stddev=0.2))
29     b = tf.Variable(tf.zeros([num_outputs], dtype = tf.float32))
30
31     # The fully connected layer
32     x = tf.add(tf.matmul(x_tensor, W), b)
33
34     # ReLU activation function
35     out = tf.nn.relu(x)
36     return out
37
38
39 def output(x_tensor, num_outputs):
40     """
41     Apply a output layer to x_tensor using weight and bias

```

```
42 : x_tensor: A 2-D tensor where the first dimension is batch size.
43 : num_outputs: The number of output that the new tensor should be.
44 : return: A 2-D tensor where the second dimension is num_outputs.
45 """
46 shape = x_tensor.get_shape().as_list()
47
48
49 # Weights and bias
50 W = tf.Variable(tf.truncated_normal([shape[1], num_outputs], dtype=tf.float32, stddev=.2))
51 b = tf.Variable(tf.zeros([num_outputs], dtype=tf.float32))
52
53
54 # The output layer
55 out = tf.add(tf.matmul(x_tensor, W), b)
56 return out
```

```
In [ ]: 1 # building the CNN network , consisting of 3 convolution layers , 3 fully connected layers
2
3
4 def conv_net(x, keep_prob):
5     """
6     Create a convolutional neural network model
7     : x: Placeholder tensor that holds image data.
8     : keep_prob: Placeholder tensor that hold dropout keep probability.
9     : return: Tensor that represents logits
10    """
11    # 3 convolution layers with max pooling
12    # All layers with same kernel, stride and maxpooling params
13
14    out = conv2d_maxpool(x,conv_num_outputs=16,conv_ksize=(3,3),conv_strides=(1,1),pool_ksize=(2,2),pool_strides
15    out = conv2d_maxpool(out,conv_num_outputs=32,conv_ksize=(3,3),conv_strides=(1,1),pool_ksize=(2,2),pool_strid
16    out = conv2d_maxpool(out,conv_num_outputs=64,conv_ksize=(3,3),conv_strides=(1,1),pool_ksize=(2,2),pool_strid
17
18
19    out = flatten(out)
20
21    # Fully connected Layer
22
23    out = fully_conn(out, num_outputs = 64)
24    out = tf.nn.dropout(out, keep_prob)
25
26    out = fully_conn(out, num_outputs = 32)
27    out = tf.nn.dropout(out, keep_prob)
28
29    out = fully_conn(out, num_outputs = 16)
30
31    #change 2no 3rd drop out done
32    # out = tf.nn.dropout(out, keep_prob)
33
34
35    out = output(out,num_outputs=10)
36
37    return out
38
```

```
In [ ]: ▶ 1 # initializing the hyperparameters for CNN
          2
          3 epochs = 100
          4 batch_size = 128
          5 keep_probability = 0.5
          6 learning_rate = 0.001
```

```
In [ ]: ▶ 1 keep_probability
```

```
In [ ]: ▶ 1 # the output of the last Fully connected layer ,i.e logits is given to conv_net( this internally
2 # calls the softmax function . )
3
4 # Resetting the default graph , so that all the previous weights, bias, inputs, etc.. are reset
5 import tensorflow as tf
6
7 tf.reset_default_graph()
8
9 # Placeholders for Inputs-- features and Labels , first dimension is the no of images
10 x = tf.placeholder(tf.float32, shape=(None, 32, 32, 3), name='input_x')
11 y = tf.placeholder(tf.float32, shape=(None, 10), name='output_y')
12
13 # no of neurons to be ignored in the Fully connected network, so that overfitting doesnt occur
14 keep_prob = tf.placeholder(tf.float32, name='keep_prob')
15
16
17 logits = conv_net(x, keep_prob)
18
19 #change 6 retrun variable name
20 #model = tf.identity(logits, name='Logits') # Name logits Tensor, so that can be loaded from disk after training
21
22 logits = tf.identity(logits, name='logits')
23 # Loss and Optimizer
24 cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=logits, labels=y))
25
26 # change 3 Learning rate not mentioned
27 #optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)
28 optimizer = tf.train.AdamOptimizer().minimize(cost)
29
30 # Accuracy
31 correct_pred = tf.equal(tf.argmax(logits, 1), tf.argmax(y, 1))
32 accuracy = tf.reduce_mean(tf.cast(correct_pred, tf.float32), name='accuracy')
```

In []: ▶

```
1 # training the neural network
2
3 def train_neural_network(session, optimizer, keep_probability, feature_batch, label_batch):
4     session.run(optimizer,
5                   feed_dict={
6                       x: feature_batch,
7                       y: label_batch,
8                       keep_prob: keep_probability
9                   })
```

In []: ▶

```
1 # printing the stats for each batch
2
3 def print_stats(session, feature_batch, label_batch, cost, accuracy):
4     """
5     Print information about loss and validation accuracy
6     : session: Current TensorFlow session
7     : feature_batch: Batch of Numpy image data
8     : label_batch: Batch of Numpy label data
9     : cost: TensorFlow cost function
10    : accuracy: TensorFlow accuracy function
11    """
12    valid_acc = sess.run(accuracy,
13                          feed_dict={
14                              x: valid_features,
15                              y: valid_labels,
16                              keep_prob: 1.
17                          })
18    train_acc = sess.run(accuracy,
19                          feed_dict={
20                              x: feature_batch,
21                              y: label_batch,
22                              keep_prob: 1.
23                          })
24
25
26    train_cost = sess.run(cost,
27                           feed_dict={
28                               x: feature_batch,
29                               y: label_batch,
30                               keep_prob: 1.
31                           })
32
33
34    print('Cost: {:.>8.5f} Accuracy on training {:.4f} '
35          'Validation Accuracy: {:.4f}'.format(train_cost, train_acc, valid_acc))
```



```
In [ ]: ▶ 1 # function to split the features list into batches , each batch of batch_size
2
3 def batch_features_labels(features, labels, batch_size):
4     """
5     Split features and labels into batches
6     """
7     for start in range(0, len(features), batch_size):
8         end = min(start + batch_size, len(features))
9         yield features[start:end], labels[start:end]
10
11 # Loading the data from the file and get batches
12 def load_preprocess_training_batch(batch_id, batch_size):
13     """
14     Load the Preprocessed Training data and return them in batches of <batch_size> or less
15     """
16     filename = 'preprocess_batch_' + str(batch_id) + '.p'
17     features, labels = pickle.load(open(filename, mode='rb'))
18
19     # Return the training data in batches of size <batch_size> or less
20     return batch_features_labels(features, labels, batch_size)
```

In []: ▶

```
1  # Training for a single batch
2  # This function trains the neural network for each batch from the features list
3  #save_model_path = './image_classification'
4  #keep_probability = 0.6
5
6  print('Training...on a single batch')
7  with tf.Session() as sess:
8      # Initializing the variables
9      sess.run(tf.global_variables_initializer())
10
11
12     # Training cycle
13     for epoch in range(epochs):
14         # for a single batch
15         batch_i = 1
16         for batch_features, batch_labels in load_preprocess_training_batch(batch_i, batch_size):
17             train_neural_network(sess, optimizer, keep_probability, batch_features, batch_labels)
18
19         print('Epoch {:>2}, CIFAR-10 Batch {:}: '.format(epoch + 1, batch_i), end='')
20         print_stats(sess, batch_features, batch_labels, cost, accuracy)
21
```

```
In [ ]: ▶ 1 # This function trains the neural network for each batch from the features list
2 import timeit
3
4
5
6 save_model_path = './image_classification'
7
8 start = timeit.timeit()
9 print('Training...')
10 with tf.Session() as sess:
11     # Initializing the variables
12     sess.run(tf.global_variables_initializer())
13     #epochs = 10
14     # Training cycle
15     for epoch in range(epochs):
16         # Loop over all batches
17         n_batches = 5
18         for batch_i in range(1, n_batches + 1 ):
19             for batch_features, batch_labels in load_preprocess_training_batch(batch_i, batch_size):
20                 train_neural_network(sess, optimizer, keep_probability, batch_features, batch_labels)
21
22             print('Epoch {:>2}, CIFAR-10 Batch {:} : '.format(epoch + 1, batch_i), end='')
23             print_stats(sess, batch_features, batch_labels, cost, accuracy)
24
25     end = timeit.timeit()
26     tm = (end-start)*1000
27
28
29     # Save Model
30     saver = tf.train.Saver()
31     save_path = saver.save(sess, save_model_path)
```

```
In [ ]: ▶ 1 # time to train the model for all the batches and epochs
2 print("The Total time to train the model is : " ,tm," ms")
```

```

In [ ]: 1 # function to display the images
2 from sklearn.preprocessing import LabelBinarizer
3 import matplotlib.pyplot as plt
4 import random
5
6 n_samples = 4
7 top_n_predictions = 5
8 model_path = save_model_path
9
10 # Defining function to load the names of the labels of whose images in the DataSet
11 def load_label_names():
12     """    Load the label names from file    """
13     return ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
14
15
16 def display_image_predictions(features, labels, predictions):
17     n_classes = 10
18
19     label_names = load_label_names()
20     label_binarizer = LabelBinarizer()
21     label_binarizer.fit(range(n_classes))
22     label_ids = label_binarizer.inverse_transform(np.array(labels))
23
24     fig, axes = plt.subplots(nrows=4, ncols=2)
25     fig.tight_layout()
26     fig.suptitle('Softmax Predictions', fontsize=20, y=1.1)
27
28     n_predictions = 5
29     margin = 0.05
30     ind = np.arange(n_predictions)
31     width = (1. - 2. * margin) / n_predictions
32     # print(width)
33
34
35     for image_i, (feature, label_id, pred_indicies, pred_values) in enumerate(zip(features, label_ids, predictions)):
36         pred_names = [label_names[pred_i] for pred_i in pred_indicies]
37         correct_name = label_names[label_id]
38
39         # axes[image_i][0].imshow(feature)
40         #change1
41         axes[image_i][0].imshow(feature)

```

```
42 axes[image_i][0].set_title(correct_name)
43 axes[image_i][0].set_axis_off()
44
45 # print(ind + margin)
46 axes[image_i][1].barh(ind + margin, pred_values[:, -1], width)
47 axes[image_i][1].set_yticks(ind + margin)
48 axes[image_i][1].set_yticklabels(pred_names[:, -1])
49 axes[image_i][1].set_xticks([0, 0.5, 1.0])
```

```

In [ ]: 1 # Testing the Model, after executing the model
2 import random
3
4 def test_model():
5     ''' To Test the model after the model is been trained '''
6
7     test_features, test_labels = pickle.load(open('preprocess_training.p', mode='rb'))
8     loaded_graph = tf.Graph()
9
10
11     with tf.Session(graph=loaded_graph) as sess:
12         # Load model
13         loader = tf.train.import_meta_graph(model_path + '.meta')
14         loader.restore(sess, model_path)
15
16
17         # Get Tensors from Loaded model
18         loaded_x = loaded_graph.get_tensor_by_name('input_x:0')
19         loaded_y = loaded_graph.get_tensor_by_name('output_y:0')
20         loaded_keep_prob = loaded_graph.get_tensor_by_name('keep_prob:0')
21         loaded_logits = loaded_graph.get_tensor_by_name('logits:0')
22         loaded_acc = loaded_graph.get_tensor_by_name('accuracy:0')
23
24         # Get accuracy in batches for memory limitations
25         test_batch_acc_total = 0
26         test_batch_count = 0
27
28         for test_feature_batch, test_label_batch in batch_features_labels(test_features, test_labels, batch_size):
29             test_batch_acc_total += sess.run(loaded_acc, feed_dict={loaded_x: test_feature_batch,
30                                                                     loaded_y: test_label_batch,
31                                                                     loaded_keep_prob: 1.0
32                                                                     })
33             test_batch_count += 1
34
35
36         print('Testing Accuracy: {}'.format(test_batch_acc_total/test_batch_count))
37
38         # Print Random Samples
39         random_test_features, random_test_labels = tuple(zip(*random.sample(list(zip(test_features, test_labels)
40         random_test_predictions = sess.run( tf.nn.top_k(tf.nn.softmax(loaded_logits), top_n_predictions),
41                                             feed_dict={loaded_x: random_test_features,loaded_y: random_test_label

```

```
42         loaded_keep_prob: 1.0 } )
43
44     display_image_predictions(random_test_features, random_test_labels, random_test_predictions)
45
46
```

In []: ▶

```
1  # testing the model
2
3  test_model()
```